

OpenCV Object Tracking Notes

Colin Burdine

1 Existing Methods

1.1 Minimum Output Sum of Squares Error (MOSSE) Filters [1]

- *Filter-Based* trackers work by initially selecting a target that is centered within a trackign window in the first frame.
- The target is then tracked by correlating the filter over a search window in the next frame, and the location that maximizes the correlation is considered the location in the next frame.
- The principle virtue of the MOSSE Filter is that it is fast, and can perform real-time tracking at high frame rates, making it suitable for edge computing purposes.
- The correlation is computed using the *Fast Fourier Transform* (FFT; denoted by \mathcal{F}). The correlation σ is given by the element-wise complex inner product of the filter matrix h and the frame image f :

$$\sigma = \mathcal{F}^{-1}(F \odot H^*) \quad \text{where} \quad F = \mathcal{F}(f), \quad H = \mathcal{F}(h)$$

- The most computationally expensive operation in this process is the 2D FFT and inverse 2D FFT, which can be performed in $\mathcal{O}(n^2 \log(n))$ for an $n \times n$ filter.
- We shall break down the operation of MOSSE filters into some critical sections:

Preprocessing

- For simplicity, the MOSSE algorithm first reduces the image to a log grayscale. Then, the image is normalized to have pixel intensity values with mean 0 and std. dev. 1.
- Since the 2D FFT projects a rectangular region onto a torus shape, connecting the top with the bottom and the left side with the right side, we want to avoid a filter “wrapping around” the edges of the image. This is remedied by multiplying the image by a cosine window.

Applying the MOSSE Filter

Ideally, a desired MOSSE filter H satisfies the condition $G = F \odot H^*$, however, we can find the filter H that is closest to this equality in the L^2 sense by solving the optimization problem that minimizes the residue.

$$\min_{H^*} \sum_i |F_i \odot H_i^* - G_i|^2$$

Since the set of possible filters forms a Hilbert space, the solution (given by the Hilbert projection Theorem) is simply the projection of G onto the basis of F , given by:

$$H^* = \frac{\sum_i G_i \odot F_i^*}{\sum_i F_i \odot F_i^*}$$

Regularization

- A Mosse filter can fit a single image (e.g. f_1) perfectly, which is not desirable in general. Thus, we use the average of the filter trained across all of the images to produce a filter that generalizes better. This averaging is motivated by the idea of *Bootstrap Aggregation*:

$$H^* = \frac{1}{N} \sum_i \frac{G_i \odot F_i^*}{F_i \odot F_i^*}$$

- MOSSE filters tend to be unstable for a low number of images; however, we can apply regularization by replacing the denominator $(F_i \odot F_i^*)$ with $(F_i \odot F_i^*) + (\epsilon \mathbf{1})$, where ϵ is the regularization parameter and $\mathbf{1}$ is the matrix of ones.

Initializing and Updating the MOSSE Filter

1.2 Generic Object Tracking Using Regression Networks (GOTURN)

1.3 MedianFlow Tracking

1.4 Discriminative Correlation Filters with Channel and Spatial Reliability (CSRT)

1.5 Kernelized Correlation Filters (KCF)

References

- [1] David Bolme, J. Beveridge, Bruce Draper, and Yui Lui. Visual object tracking using adaptive correlation filters. pages 2544–2550, 06 2010.